

# Databases

- Software that stores data on disk
- Runs as a server and is communicated with via TCP sockets
- Provides an API to store/retrieve data
  - The software handles the low-level file IO
  - Allows us to think about our data, not how to store it
- Provides many optimizations

# Databases

- We'll look at 2 different databases
- Both are pieces of software that must be downloaded, installed, ran, then connected to via TCP
- MySQL
  - A server implementing SQL (Structured Query Language)
- MongoDB
  - Am unstructured server based on document stores

# MySQL

- Once you download, install, and run the server
  - It will listen for TCP connections on port 3306 (By default)
- Install a library for your language that will connect to the MySQL server
  - You will not have to connect to your database at the TCP level in this course (True for MongoDB as well)
- The library will provide a convenient API
  - Send queries using a query language

# MySQL - Connection

- MySQL runs and you install a library to connect to it
- Connect to MySQL Server by providing:
  - The url of the database
  - username/password for the database
  - Whatever you chose when setting up the database

```
val url = "jdbc:mysql://localhost/mysql"  
val username = "root"  
val password = "12345678"
```

```
var connection: Connection = DriverManager.getConnection(url, username, password)
```



# MySQL - Insert Data

- Once connected, we can send SQL statements to the server

```
val statement = connection.createStatement()  
statement.execute("CREATE TABLE IF NOT EXISTS players (username TEXT, points INT)")
```

- If using inputs from the user, always use prepared statements

```
val statement = connection.prepareStatement("INSERT INTO players VALUE (?, ?)")  
  
statement.setString(1, "mario")  
statement.setInt(2, 10)  
  
statement.execute()
```

# MySQL - Security

- Not using prepared statements?
  - **Vulnerable to SQL injection attacks**
- If you concatenate user inputs directly into your SQL statements
  - Attacker chooses a username of `'';DROP TABLE players;`
  - You lose all your data
  - Even worse, they find a way to access the entire database and steal other users' data
  - SQL Injection is the most common successful attack on servers

# MySQL - Retrieve Data

- Send queries to pull data from the database

```
val statement = connection.createStatement()
val result: ResultSet = statement.executeQuery("SELECT * FROM players")

var allScores: Map[String, Int] = Map()

while (result.next()) {
    val username = result.getString("username")
    val score = result.getInt("points")
    allScores = allScores + (username -> score)
}
```

# SQL

- SQL is based on tables with rows and column
  - Similar in structure to CSV except the values have types other than string
- How do we store an array or key-value store?
  - With CSV our answer was to move on to JSON
  - SQL answer is to create a separate table and use JOINS
  - Or, try MongoDB



# MongoDB

- Runs on port 27017 (By default)
- A document-based database
- Instead of using tables, stores data in a structure very similar to JSON
- In python/JS
  - Insert dictionaries/objects directly
- Each object is stored in a collection

# MongoDB - Connection

- Download a connection library and use to establish a connection with MongoDB
- MongoDB is separated into several layers
  - Databases - Named by Strings; Contains collections
  - Collections - Where the data is stored; similar to a SQL table
- Access your collections to insert/retrieve/update/delete data

```
from pymongo import MongoClient

mongo_client = MongoClient("localhost")
db = mongo_client["cse312"]
chat_collection = db["chat"]
```

# MongoDB - Insert Data

- Insert dictionaries/objects directly
- For languages without a data structure comparable to dictionaries/objects
- More work to do to prepare your data for Mongo

```
chat_collection.insert_one({"username": "hartloff", "message": "hello"})
```

# MongoDB - Security

- No Mongo injection attacks
- Mongo does not rely on parsing statements like SQL
- Any injected code would be treated as values
- It's like using prepared statements all the time with no extra work!

```
chat_collection.insert_one( {"username": "hartloff", "message": "hello"})
```



# MongoDB - Retrieve Data

- Retrieve documents using find
- Find takes a key-value store and returns all documents with those values stored at the given keys
  - Ex. {"username": "hartloff"} returns all documents with a username of "hartloff"
- To retrieve all documents, use an empty key-value store {}

```
my_data = chat_collection.find({"username": "hartloff"})  
all_data = chat_collection.find({})
```

# MongoDB vs. SQL

- MongoDB is unstructured
  - Can add objects in any format to a collection
  - Can mix formats in a single collection
    - I.e. In a single collection the documents can have different attributes
- SQL is structured (That's what the S stands for)
  - Table columns must be pre-defined
    - All rows have the same attributes
    - Adding a column can be difficult
  - Fast!

# MongoDB vs. SQL

- Hot Take
  - MongoDB is best for prototyping when the structure of your data is constantly changing
    - Take advantage of the flexibility
  - SQL is best once your data has a defined structure
    - Take advantage of the efficiency

# Databases in CSE312

- You're expected to find documentation/tutorials for your database and language of choice
- Choose a database
- Find a connection library for that database in your language
- Add the library to your dependancies
  - Make sure you install it in your Dockerfile
- Study documentation to learn how to use the database



# Docker Compose Revisited

docker-compose.yml

```
version: '3.3'
services:
  mongo:
    image: mongo:4.2.5
  app:
    build: .
    environment:
      WAIT_HOSTS: mongo:27017
    ports:
      - '8080:8000'
```

- Let's modify our docker compose configuration to run our database

# Docker Compose

docker-compose.yml

```
version: '3.3'
services:
  mongo:
    image: mongo:4.2.5
  app:
    build: .
    environment:
      WAIT_HOSTS: mongo:27017
    ports:
      - '8080:8000'
```

- "services" is a list of all the images/containers to create
- We'll add a second service for the DB

# Docker Compose

`docker-compose.yml`

```
version: '3.3'
services:
  mongo:
    image: mongo:4.2.5
  app:
    build: .
    environment:
      WAIT_HOSTS: mongo:27017
    ports:
      - '8080:8000'
```

- Name each service
- These names are used as the hostnames for each container
- Used to communicate between containers

# Docker Compose

`docker-compose.yml`

```
version: '3.3'
services:
  mongo:
    image: mongo:4.2.5
  app:
    build: .
    environment:
      WAIT_HOSTS: mongo:27017
    ports:
      - '8080:8000'
```

- This service named 'mongo' uses a pre-build image
- Same as having a 1-line Dockerfile:
  - "FROM mongo:4.2.5"
- No Dockerfile is needed



# Docker Compose

docker-compose.yml

```
version: '3.3'
services:
  mongo:
    image: mongo:4.2.5
  app:
    build: .
    environment:
      WAIT_HOSTS: mongo:27017
    ports:
      - '8080:8000'
```

- Use 'environment' to set any needed environment variables
- If using MySQL, set variables for your username/password

# Docker Compose

## docker-compose.yml

```
version: '3.3'
services:
  mongo:
    image: mongo:4.2.5
  app:
    build: .
    environment:
      WAIT_HOSTS: mongo:27017
    ports:
      - '8080:8000'
```

- We use an environment variable to tell our app to wait until the database is running before connecting to it

```
FROM python:3.8.2
```

```
ENV HOME /root
```

```
WORKDIR /root
```

```
COPY . .
```

```
RUN pip install -r requirements.txt
```

```
EXPOSE 8000
```

```
ADD https://github.com/ufoscout/docker-compose-wait/releases/download/2.2.1/wait /wait
```

```
RUN chmod +x /wait
```

```
CMD /wait && python app.py
```

# Docker Compose

## docker-compose.yml

```
version: '3.3'
services:
  mongo:
    image: mongo:4.2.5
  app:
    build: .
    environment:
      WAIT_HOSTS: mongo:27017
    ports:
      - '8080:8000'
```

- If the app runs before the database, it won't be able to establish a DB connection
- Solution: Wait for the DB to start before running the app

```
FROM python:3.8.2
```

```
ENV HOME /root
```

```
WORKDIR /root
```

```
COPY . .
```

```
RUN pip install -r requirements.txt
```

```
EXPOSE 8000
```

```
ADD https://github.com/ufoscout/docker-compose-wait/releases/download/2.2.1/wait /wait
```

```
RUN chmod +x /wait
```

```
CMD /wait && python app.py
```

# Docker Compose

## docker-compose.yml

```
version: '3.3'
services:
  mongo:
    image: mongo:4.2.5
  app:
    build: .
    environment:
      WAIT_HOSTS: mongo:27017
    ports:
      - '8080:8000'
```

- This solution from github user "ufoscout" works well

```
FROM python:3.8.2
```

```
ENV HOME /root
```

```
WORKDIR /root
```

```
COPY . .
```

```
RUN pip install -r requirements.txt
```

```
EXPOSE 8000
```

```
ADD https://github.com/ufoscout/docker-compose-wait/releases/download/2.2.1/wait /wait
```

```
RUN chmod +x /wait
```

```
CMD /wait && python app.py
```



# Docker Compose

`docker-compose.yml`

```
version: '3.3'
services:
  mongo:
    image: mongo:4.2.5
  app:
    build: .
    environment:
      WAIT_HOSTS: mongo:27017
    ports:
      - '8080:8000'
```

- This file is used to build both images and run both containers using docker-compose

# Docker Compose

docker-compose.yml

```
version: '3.3'
services:
  mongo:
    image: mongo:4.2.5
  app:
    build: .
    environment:
      WAIT_HOSTS: mongo:27017
    ports:
      - '8080:8000'
```

~~mongo\_client = MongoClient('localhost')~~

mongo\_client = MongoClient('mongo')

- Recall that we chose names for each service
- When connecting to the database in your app
  - The service name is the hostname for the container

# Docker Compose

docker-compose.yml

```
version: '3.3'
services:
  mongo:
    image: mongo:4.2.5
  app:
    build: .
    environment:
      WAIT_HOSTS: mongo:27017
    ports:
      - '8080:8000'
```

```
mongo_client = MongoClient('localhost')
```

```
mongo_client = MongoClient('mongo')
```

- Use the name of the service
- docker-compose will resolve this hostname to the appropriate container

# Docker Compose

docker-compose.yml

```
version: '3.3'
services:
  mysupercooldatabase:
    image: mongo:4.2.5
  app:
    build: .
    environment:
      WAIT_HOSTS: mongo:27017
    ports:
      - '8080:8000'
```

```
mongo_client = MongoClient('localhost')
```

```
mongo_client = MongoClient('mysupercooldatabase')
```

- We can name our services whatever we want
- Make sure you are consistent!



# Running Your App

- `docker-compose up --build --force-recreate`
- Will now start both containers
- Use the service name as the host name to communicate across containers



# HTML Templates



# The Problem

- We want to serve custom HTML
- You want to build a chat feature for your app
  - Users will submit their messages
  - Messages will appear to all users
  - Messages are contained in your HTML
- How do we serve HTML that will change as users send messages?



# HTML Templates

- Instead of writing complete HTML files
  - Write HTML templates
- An HTML template is an "incomplete" HTML file that is used to generate complete pages
- Use additional markup to add placeholders in the HTML
- Replace the placeholders with data at runtime



# HTML Templates

- Example template with 3 placeholders
- The title, description, and image\_filename will be replaced later
- Provide values for these 3 placeholders to serve a response

```
<!DOCTYPE html>
<html lang="en">
<head>
  <meta charset="UTF-8">
  <title>This page was generated from a template</title>
</head>
<body>

<h1>{{title}}</h1>

<p>{{description}}</p>



</body>
</html>
```



# HTML Templates

- To substitute the placeholders
  - Use any string manipulation that gets the job done
  - Find/replace is the simplest solution
  - May want more advanced approaches if you want to add more functionality

```
<!DOCTYPE html>
<html lang="en">
<head>
  <meta charset="UTF-8">
  <title>This page was generated from a template</title>
</head>
<body>

<h1>{{title}}</h1>

<p>{{description}}</p>



</body>
</html>
```



# Common Template Features

- Loops
- To add loops to your templates
  - Choose syntax for the start and end of the loop

```
<!DOCTYPE html>
<html lang="en">
<head>
  <meta charset="UTF-8">
  <title>This page was generated from a template</title>
</head>
<body>

  {{loop}}
  <h6>{{content_from_data_structure}}</h6>
  {{end_loop}}

</body>
</html>
```



# Common Template Features

- Use string manipulation to find the start and end tags
- Iterate over your data
- Add the contained HTML with the placeholder replaced for each value of your data

```
<!DOCTYPE html>
<html lang="en">
<head>
  <meta charset="UTF-8">
  <title>This page was generated from a template</title>
</head>
<body>

  {{loop}}
  <h6>{{content_from_data_structure}}</h6>
  {{end_loop}}

</body>
</html>
```



# Common Template Features

- Conditionals
- Can use similar approach as loops
- Choose syntax for the start and end of each block in the conditional

```
<!DOCTYPE html>
<html lang="en">
<head>
  <meta charset="UTF-8">
  <title>This page was generated from a template</title>
</head>
<body>

  {{if cookie_set}}
  <h6>Welcome back!</h6>
  {{else}}
  <h6>Welcome!</h6>
  {{end_if}}

</body>
</html>
```



# Common Template Features

- Search for your tags
- Extract and evaluate the conditional
  - Choose how this will be evaluated
- Add the appropriate block of HTML to the page

```
<!DOCTYPE html>
<html lang="en">
<head>
  <meta charset="UTF-8">
  <title>This page was generated from a template</title>
</head>
<body>

  {{if cookie_set}}
  <h6>Welcome back!</h6>
  {{else}}
  <h6>Welcome!</h6>
  {{end_if}}

</body>
</html>
```