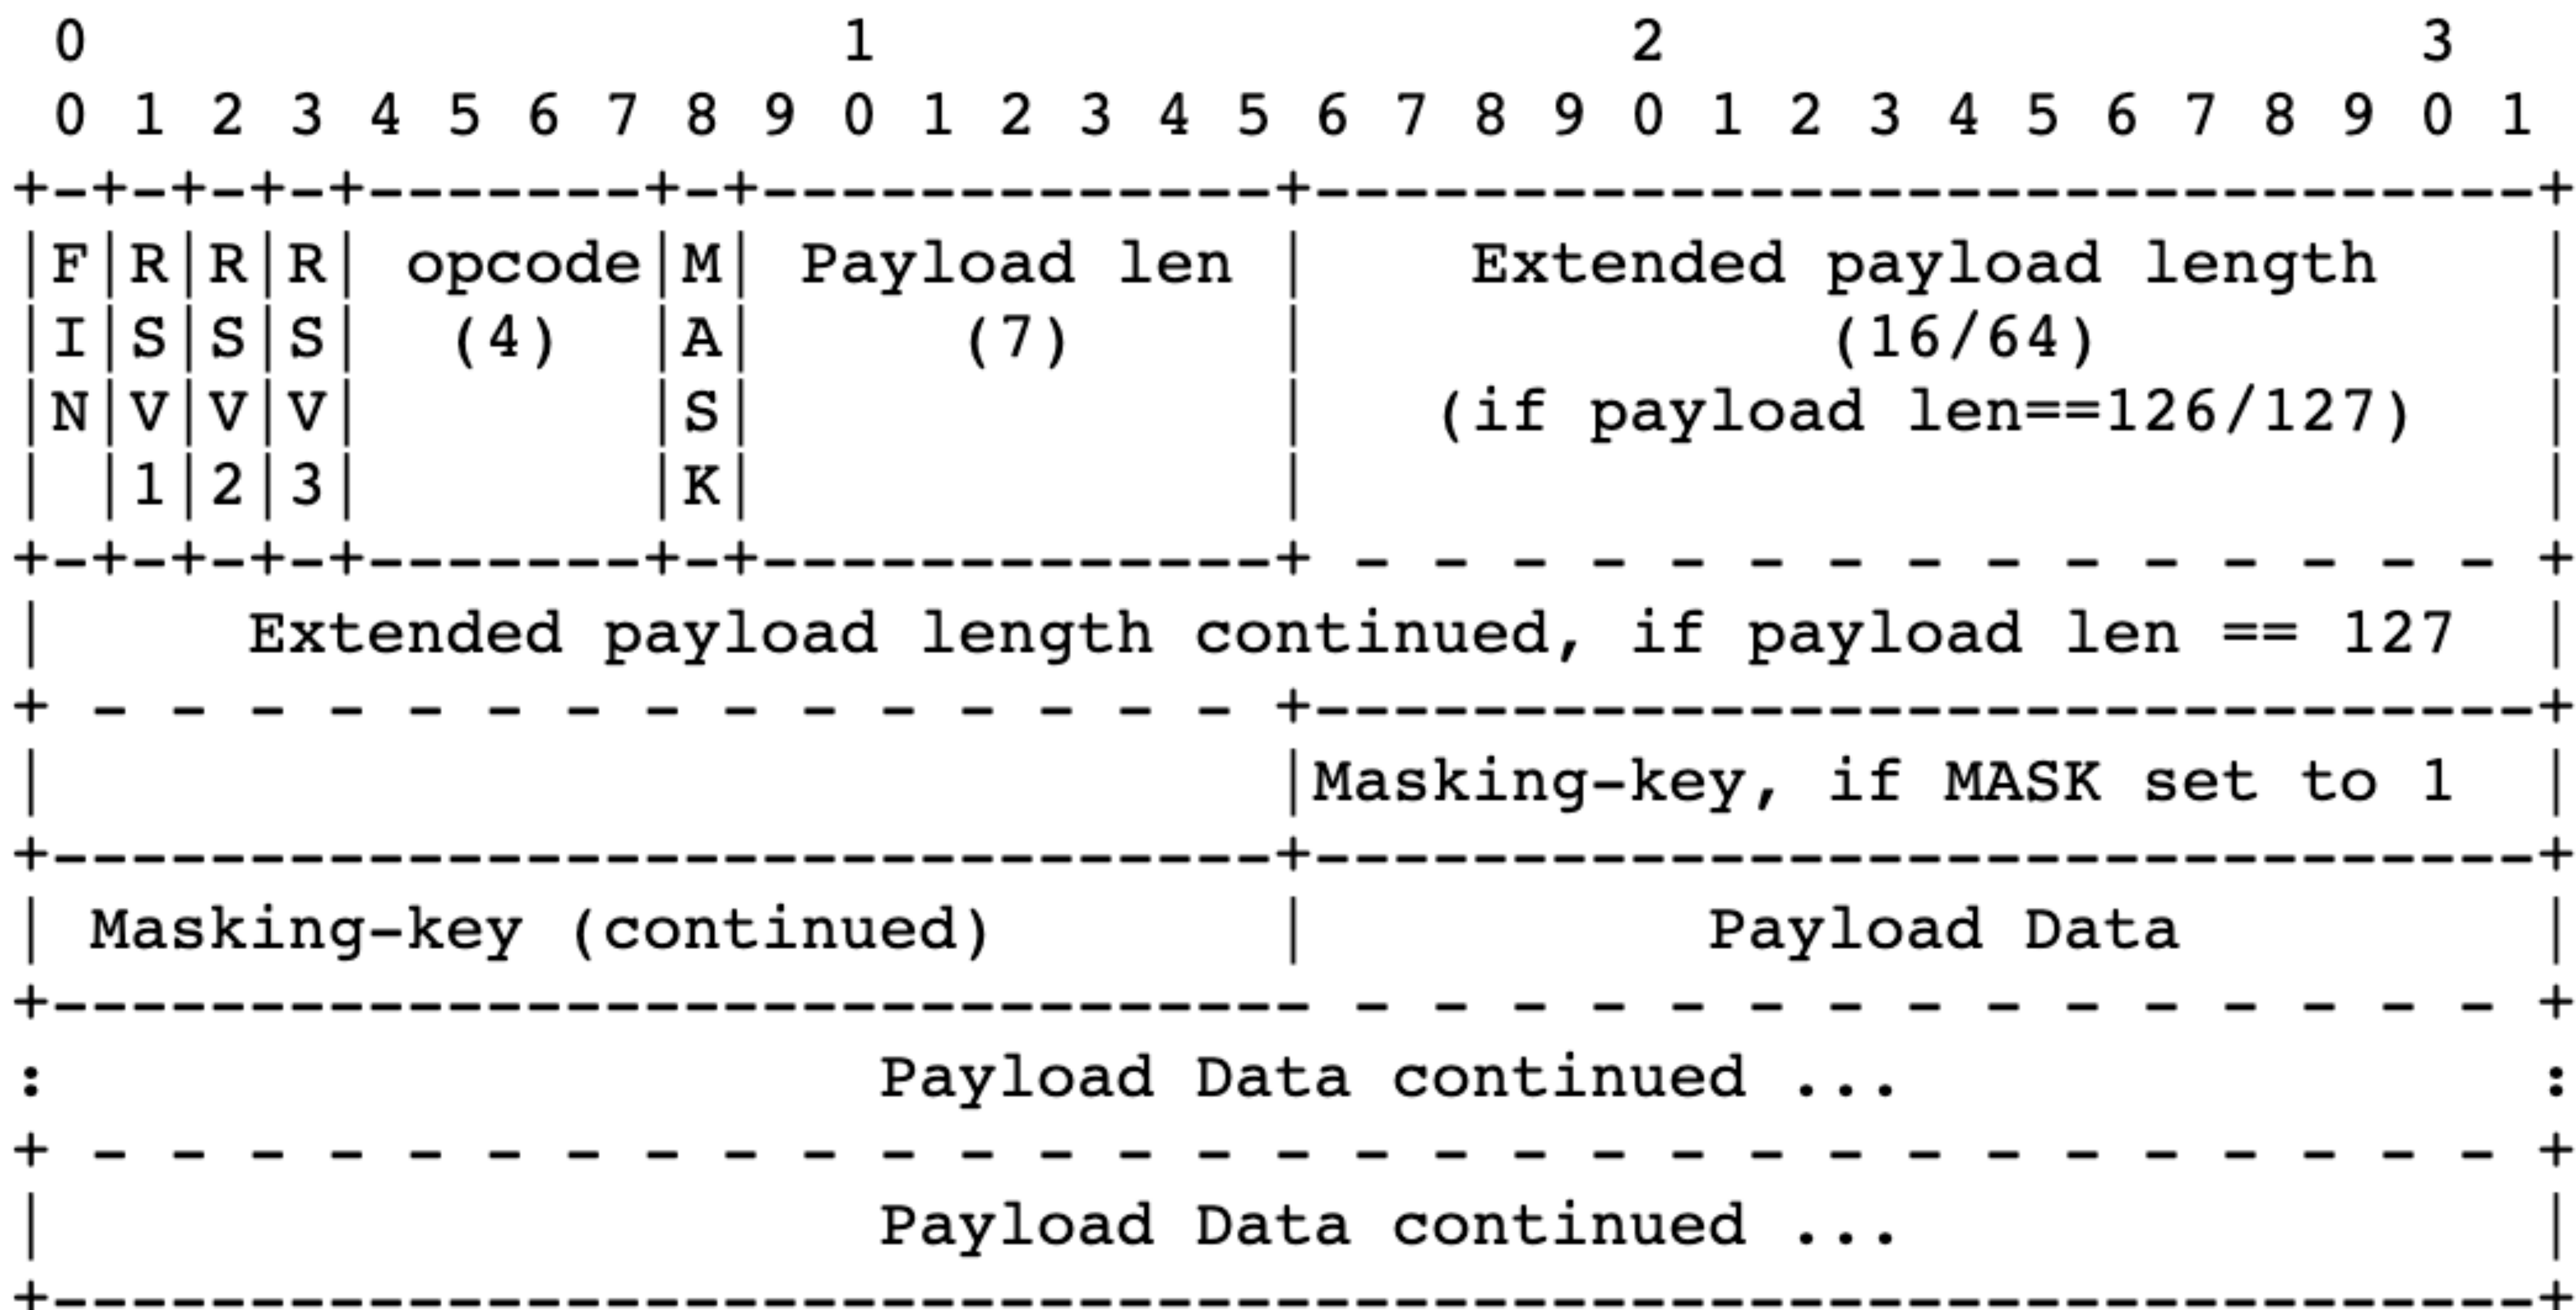


WebSocket Buffers



Special Cases

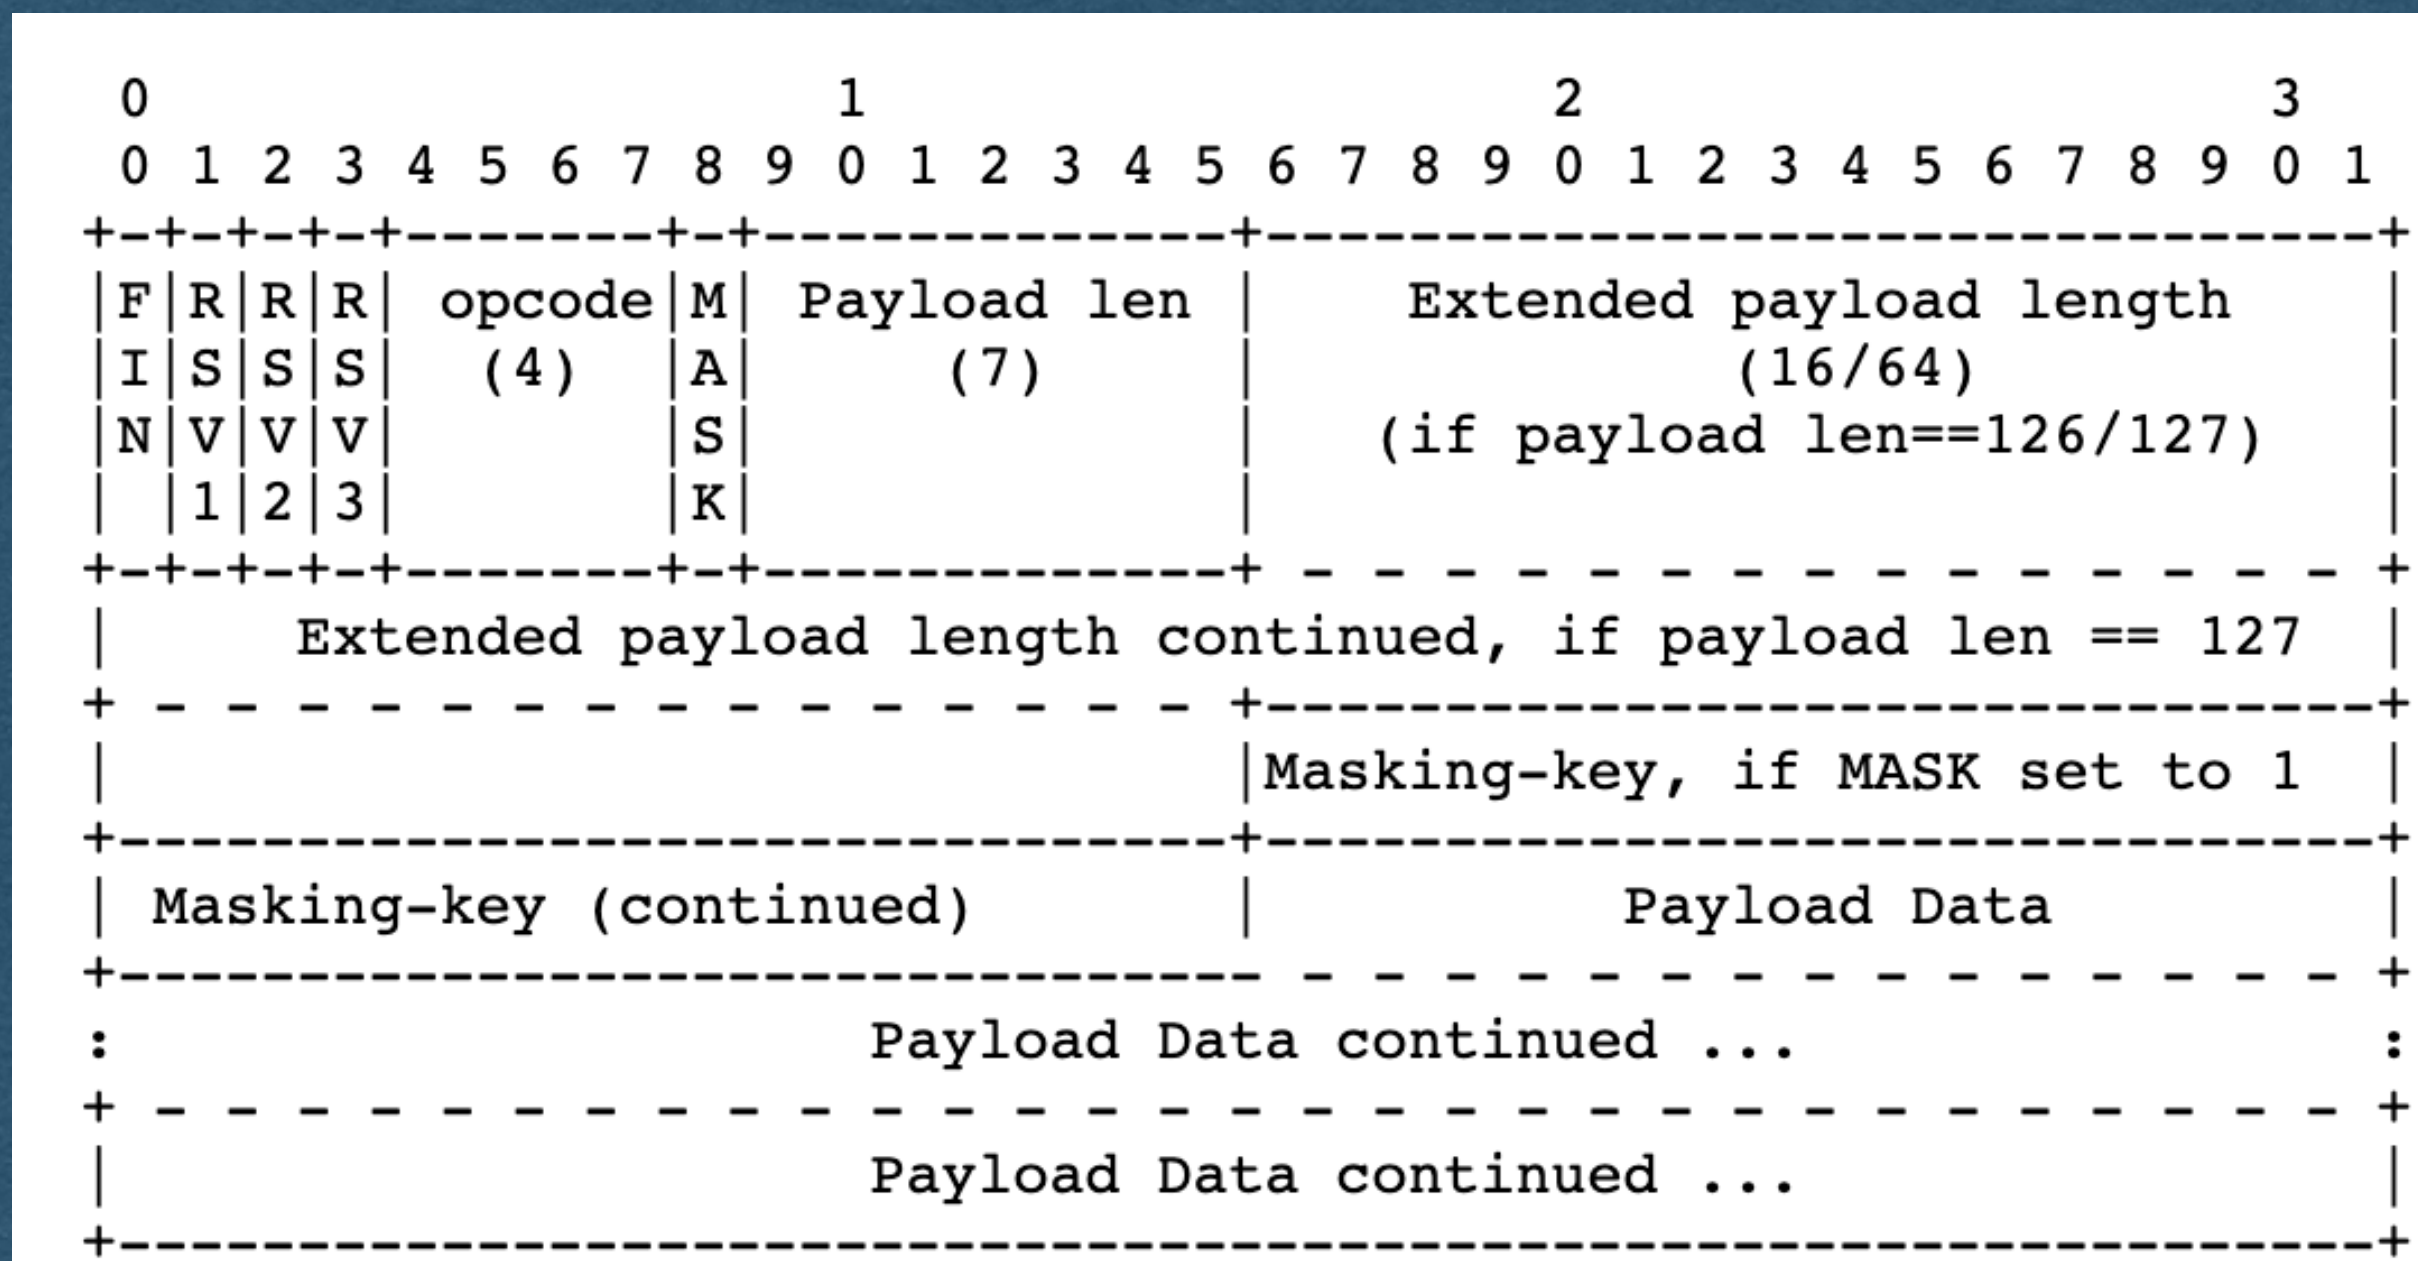
- Let's talk about 3 special cases that will come up when implementing WebSockets
 - Buffering large frames
 - Multiple frames per message using the fin bit and continuation frames
 - Multiple frames being sent back-to-back messages

Buffering Large Frames

- You will sometimes receive WebSocket frames that are large enough that they need to be buffered
 - Buffering frames is very similar to buffering HTTP requests
- When receiving a WebSocket Frame:
 - Read bytes from the socket
 - Parse the headers
 - Read the payload length from the headers
 - Keep reading bytes from the socket until you've read the entire frame
 - Payload length does not include the header bytes
 - Process the request

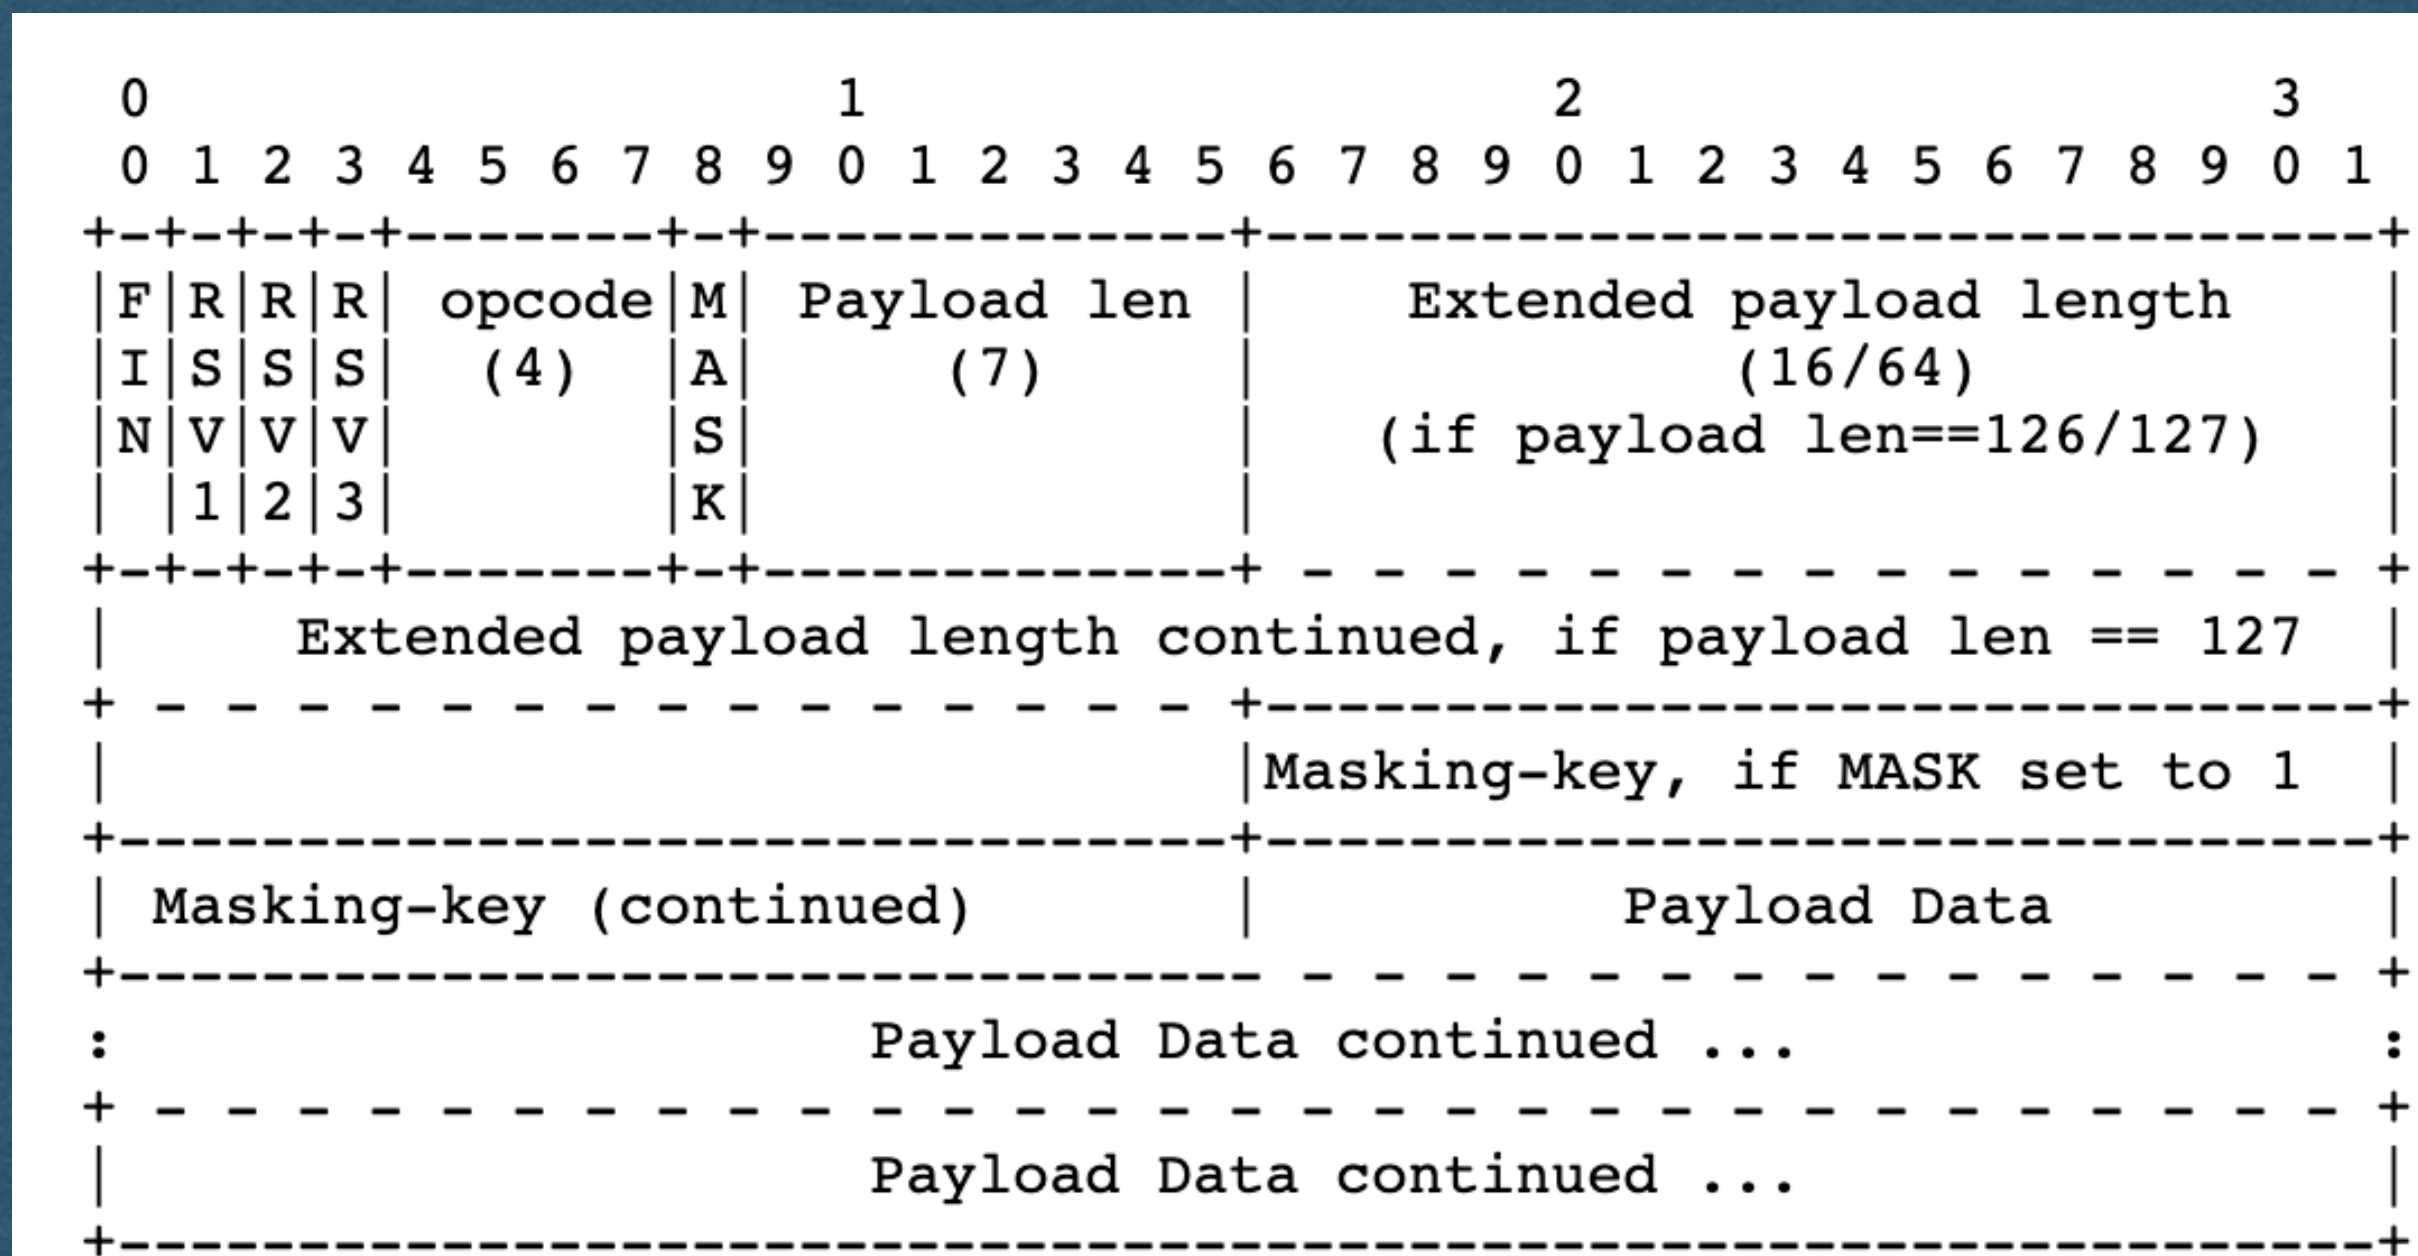
Continuation Frames

- You will sometimes receive very large messages from client that will be sent in multiple frames (>131,000 bytes in Chrome)
- Fin bit will be 0 until the last frame
- opcode will be 0000 for all but the first frame
- Payload length is **only** the length of that frame



Continuation Frames

- When you read a frame with a fin bit of 0:
 - Keep reading frames until you read a frame with a fin bit of 1
 - Combine the payload of all frames, then process the entire message



Continuation Frames

- Example of one message sent over 3 frames

Frame 1

fin bit: 0
payload length: 11
payload: "Never look "

Frame 2

fin bit: 0
payload length: 14
payload: "directly into "

Frame 3

fin bit: 1
payload length: 7
payload: "the sun"

Message: Never look directly into the sun

Back-to-back Frames

- Multiple WebSocket frames can be sent back-to-back on the same connection
 - Especially when continuation frames are used
- If you read more bytes than you expect, you have read the headers of the next frame
- Use the payload length to know how many bytes to expect
 - If you read $<$ payload length bytes, you should buffer
 - If you read $>$ payload length bytes, store the extra bytes as the start of the next frame

Back-to-back Frames

- To test for back-to-back frames:
 - Edit the front end JavaScript to send a message multiple times when the user sends a message
 - Make sure each message is duplicated the correct number of times

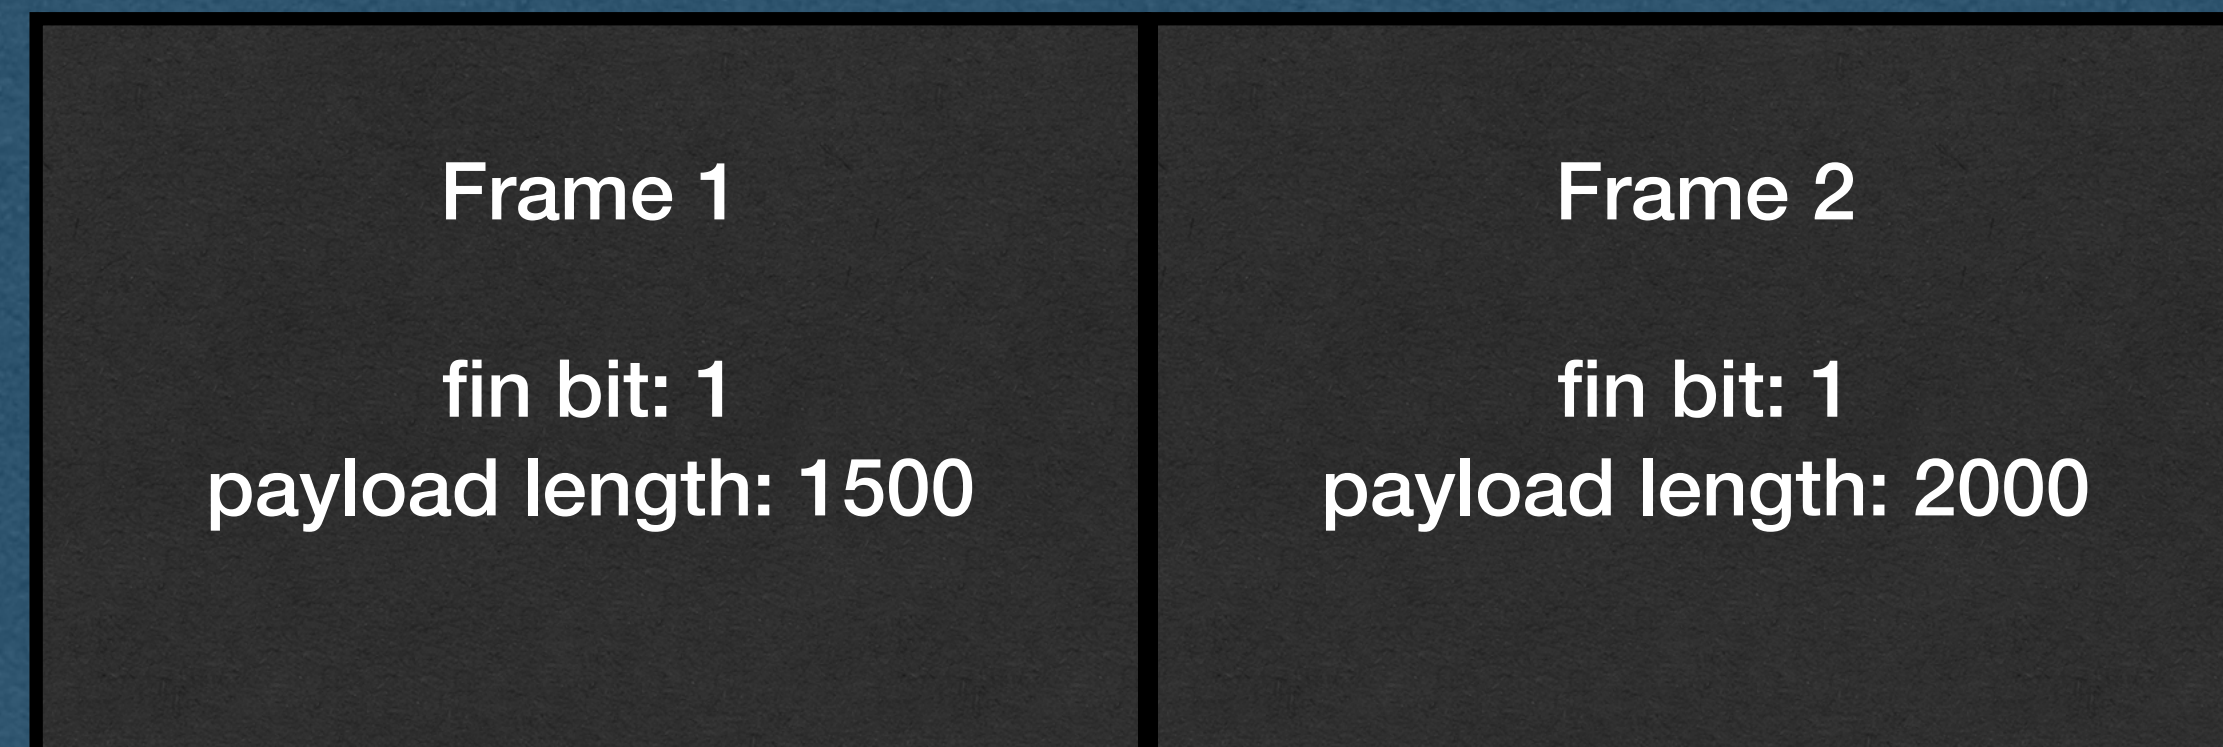
```
socket.send(JSON.stringify(value: {'messageType': 'chatMessage', 'message': message}));  
socket.send(JSON.stringify(value: {'messageType': 'chatMessage', 'message': message}));  
socket.send(JSON.stringify(value: {'messageType': 'chatMessage', 'message': message}));  
socket.send(JSON.stringify(value: {'messageType': 'chatMessage', 'message': message}));  
socket.send(JSON.stringify(value: {'messageType': 'chatMessage', 'message': message}));
```

- With this modification, every sent message should appear in chat 5 times

Back-to-back Frames

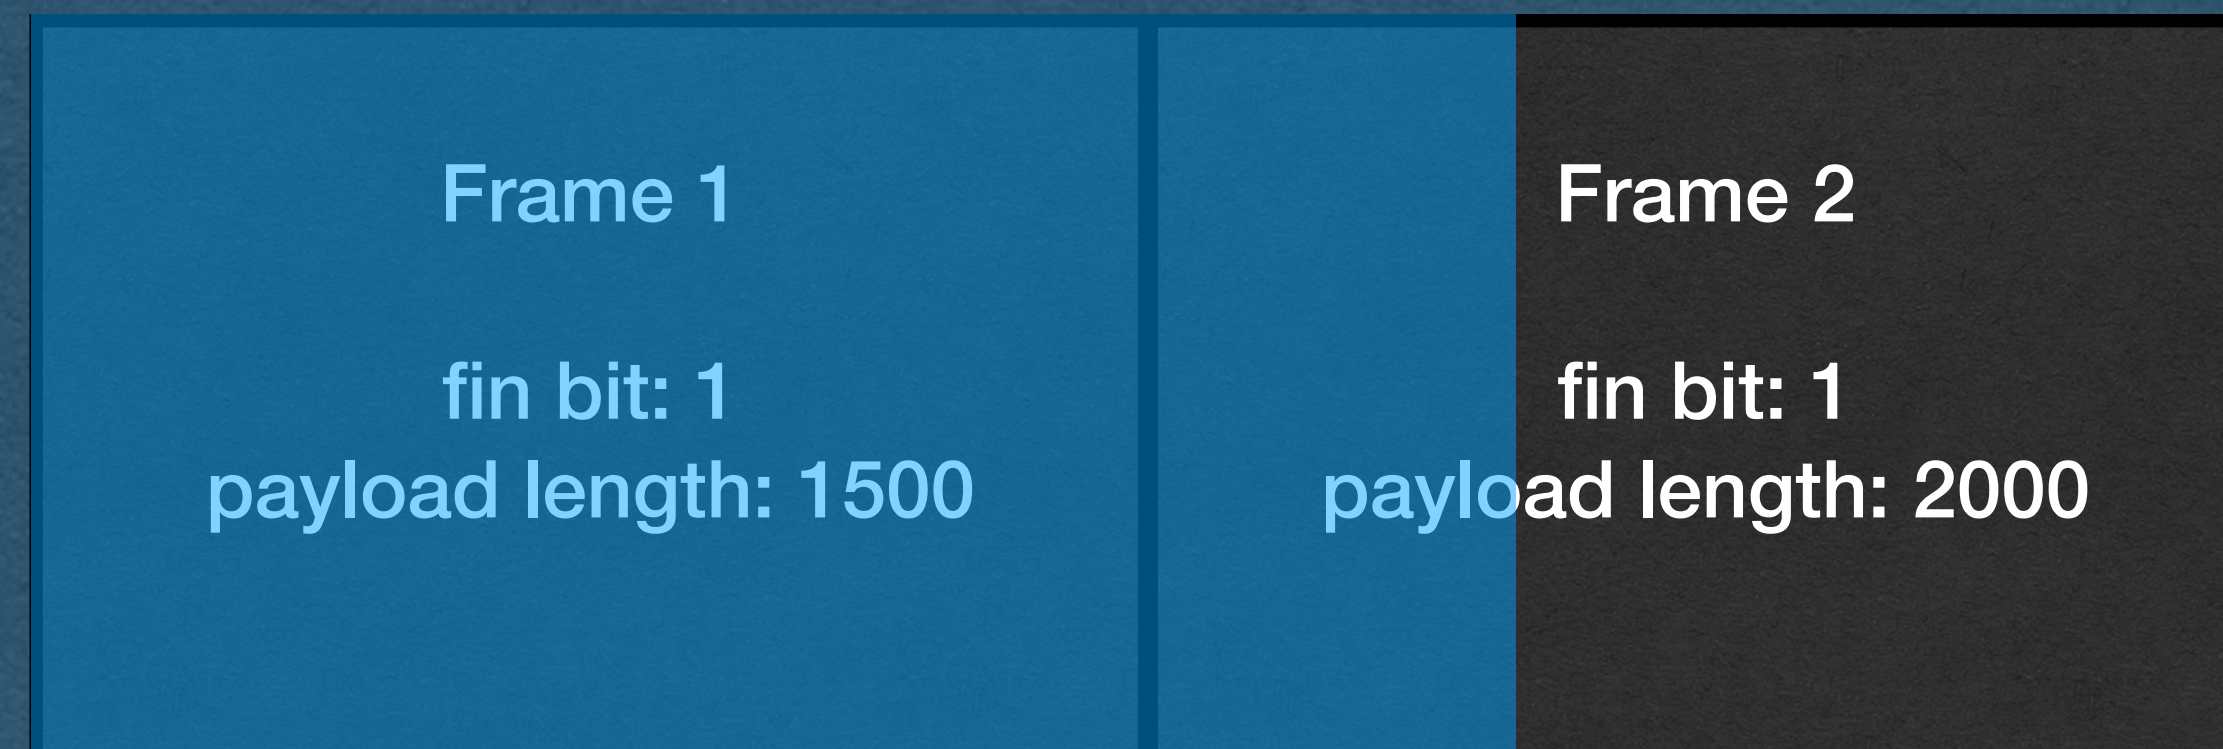
Example

- Two frames are sent by the client back-to-back
 - The first frame has a payload length of 1500
 - 1508 total bytes including headers
 - The first frame has a payload length of 2000
 - 2008 total bytes including headers
- Both frames have been processed by TCP and are ready to be read



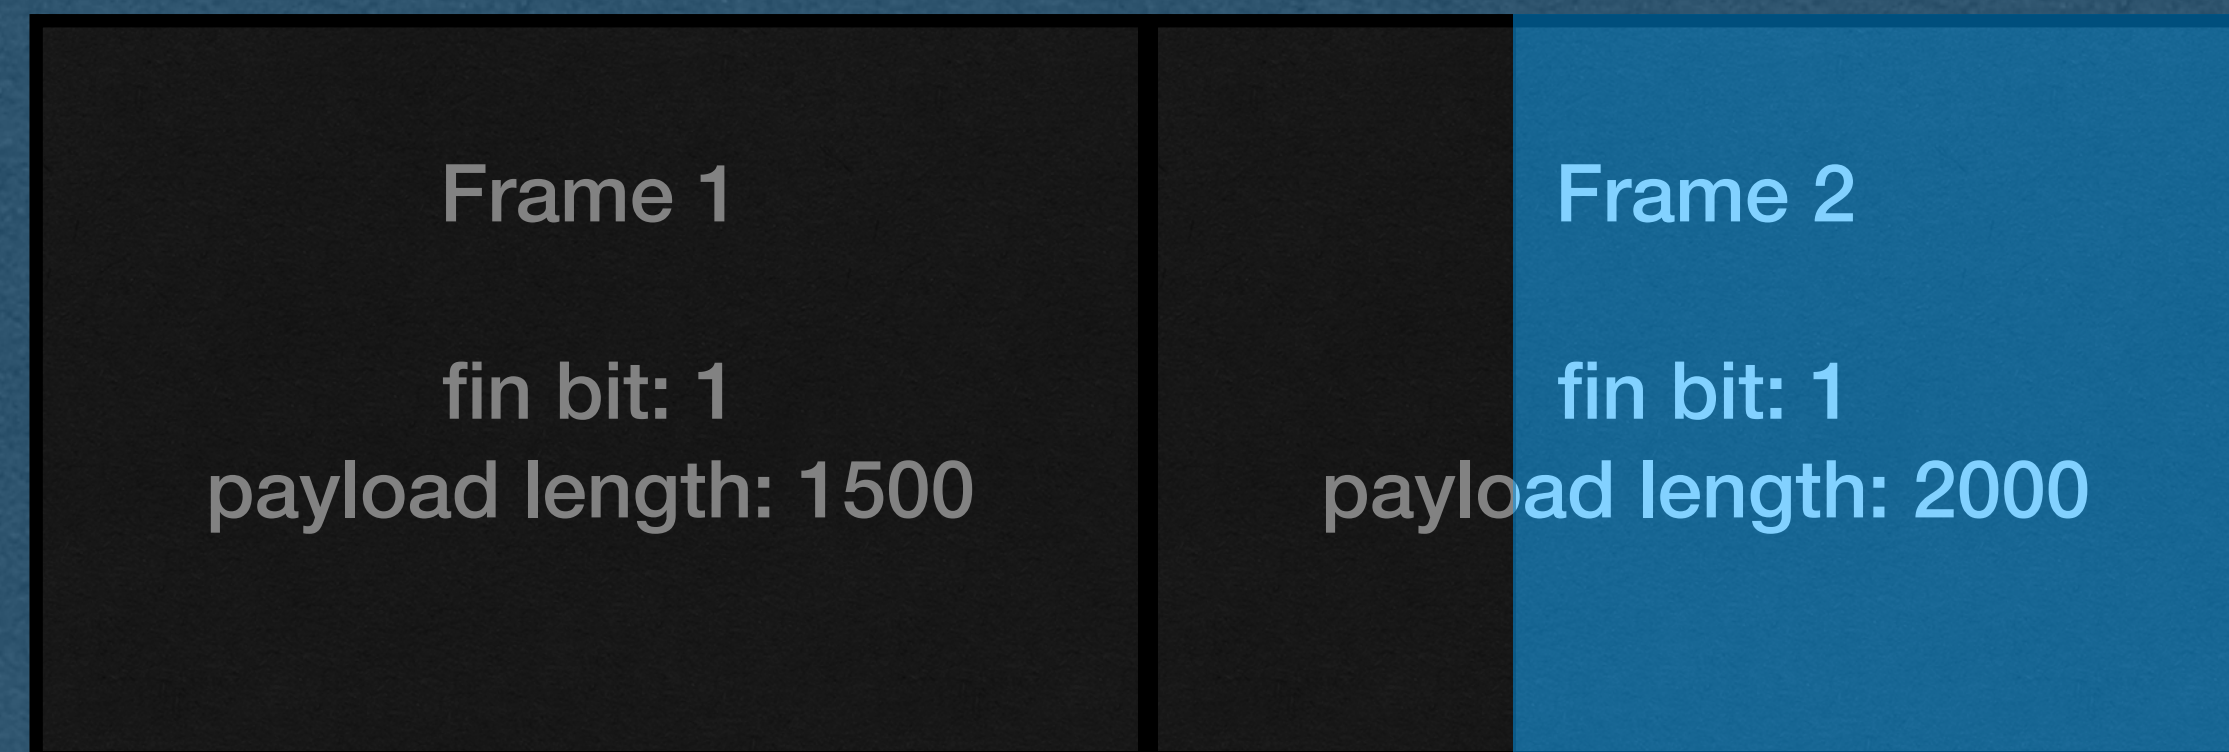
Back-to-back Frames

- There are 3516 bytes ready to be read from the socket
- And you call `received_data = self.request.recv(2048)`
 - This will read unto 2048 bytes from the socket
 - There are >2048 bytes ready, so you get the first 2048
- You read the entirety of the first frame, and the first 540 bytes of the second frame



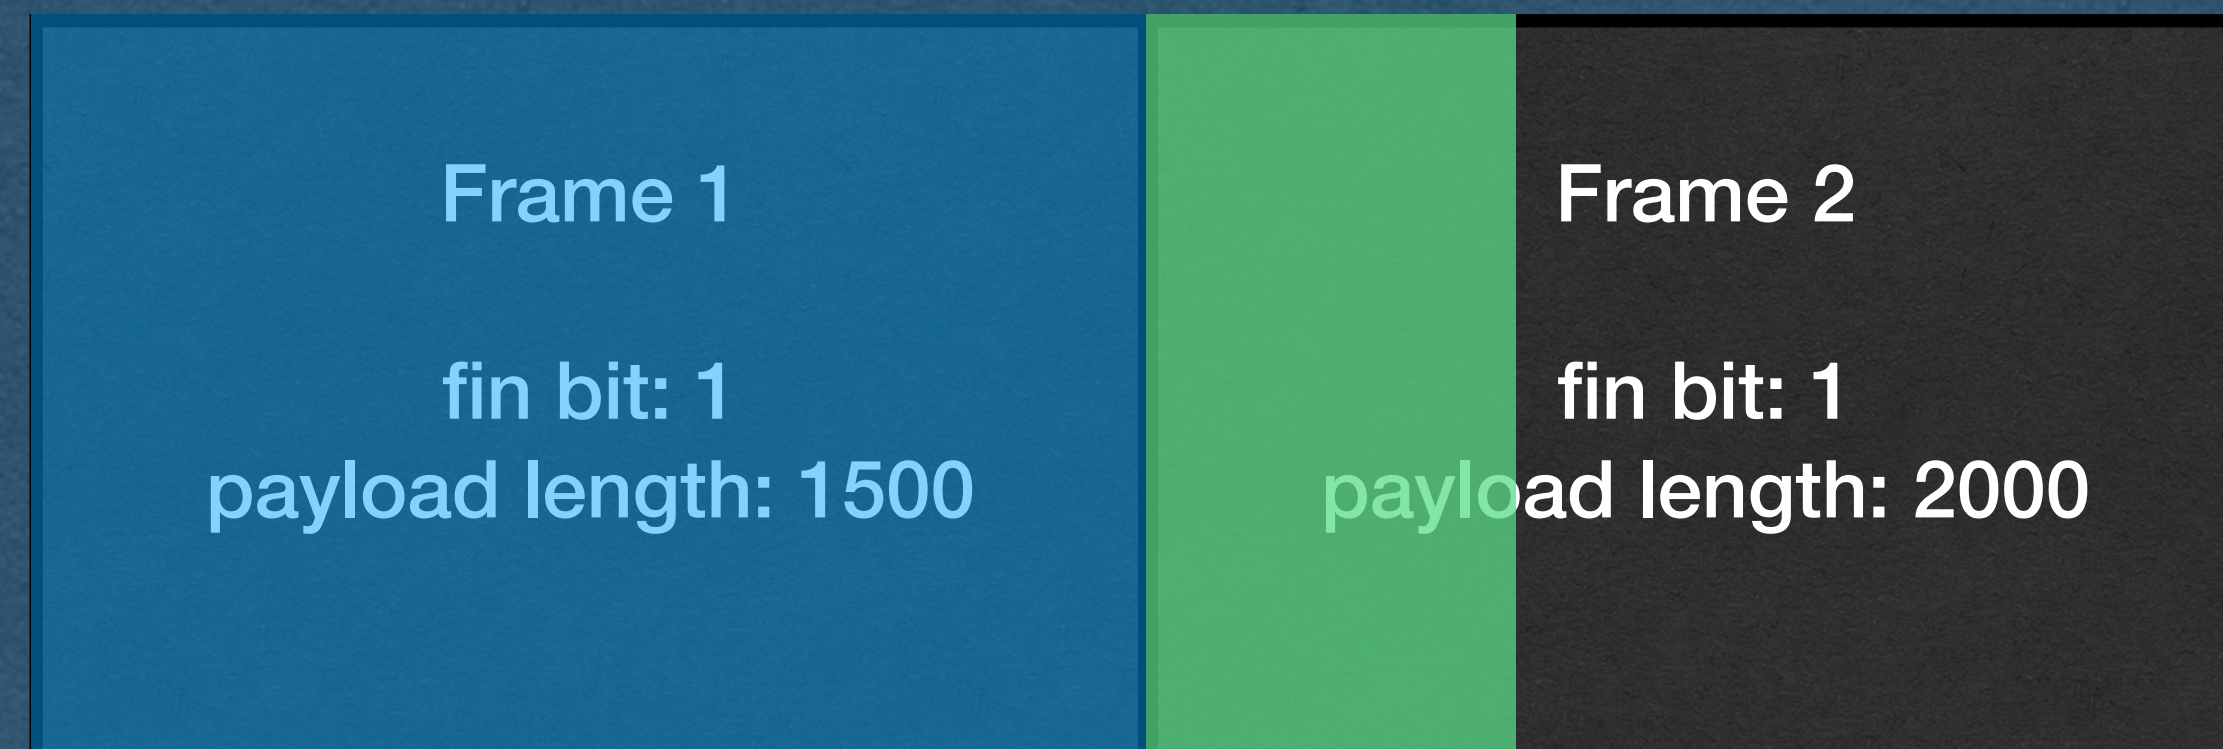
Back-to-back Frames

- If you're not careful, your loop will go back to the socket and read the remaining 1468 bytes of the second frame and attempt to parse it
- Since you start in the middle of the frame, you will run header parsing code on masked payload bytes
- You will get errors!



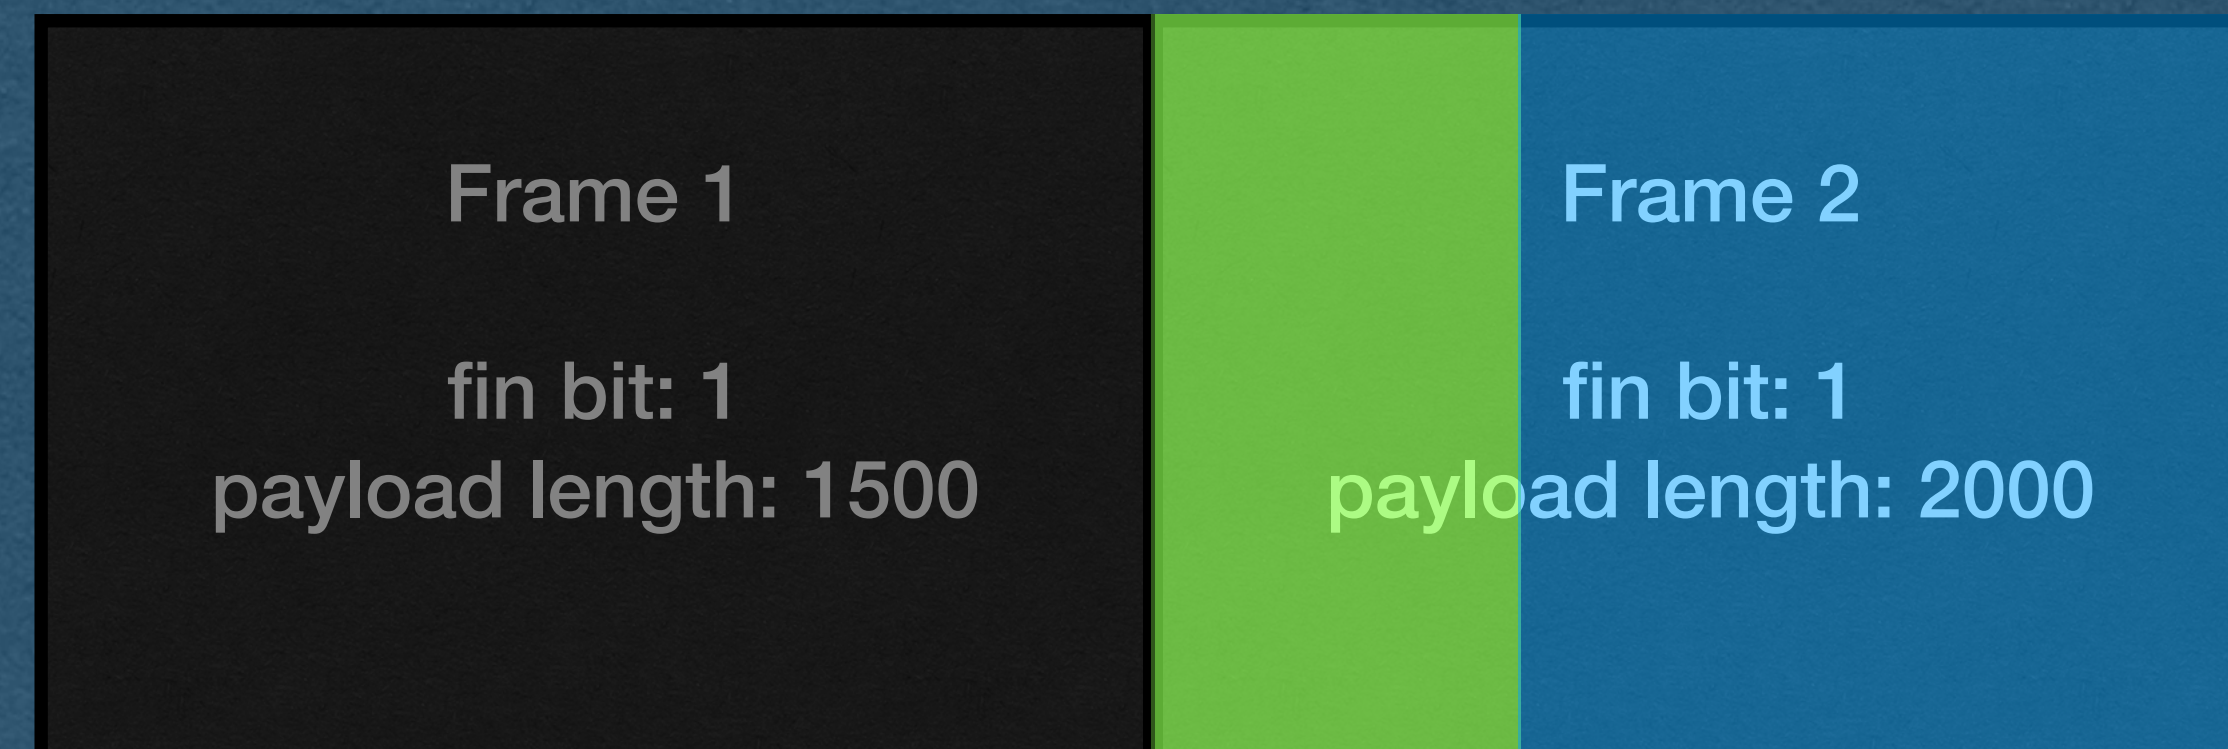
Back-to-back Frames

- When parsing the first frame:
 - Use the payload length to detect that you've read too many bytes
 - Store the extra bytes in a separate variable
 - Parse the first frame



Back-to-back Frames

- When you finish processing the first frame, start parsing the second frame with the bytes stored in the operate variable
- Check the payload length and buffer if needed to read the rest of the frame
- Recommendation: Use your top-level loop to do this so you can handle any number of back-to-back frames



Demos