

HTTP - Request/ Response

HTTP - Documentation

- HTTP/1.1 is defined by RFC2616 of the IETF
 - <https://tools.ietf.org/html/rfc2616>
 - This is THE document for all your questions about HTTP
 - Today we'll discuss topics in sections 4, 5, and 6
- RFC
 - Request For Comments
 - Submit an RFC for public discussion or to publish information
- IETF
 - Internet Engineering Task Force
 - Adopts some RFC's as Internet standards

New Lines

- A new line character in an HTTP request must be:
 - `"\r\n"`
 - Carriage return (From the days of typewriters)
 - New line
 - In the documentation this is referred to as a CRLF
 - CRLF == Carriage Return Line Feed
- Also from the documentation

The line terminator for message-header fields is the sequence CRLF. However, we recommend that applications, when parsing such headers, recognize a single LF as a line terminator and ignore the leading CR.

Request

- We'll use this simple request as an example
- What specific resource is being requested?

GET / HTTP/1.1
Host: cse312.com

The Request Line

- The first line of the request is always the request line
- The request line has 3 values separated by spaces
 - The request type (ex. GET/POST)
 - The path of the request (ex. "/")
 - The HTTP Version
 - We'll always use HTTP/1.1 in this course
 - You can assume the request uses HTTP/1.1 in your assignments without checking this string

GET / HTTP/1.1

Host: cse312.com

The Request Line

- Parse the request line by looking for the 2 space characters
 - Separate the values and check the strings
- Typically: When the root path "/" is requested, serve your home page
 - By convention, web servers look for index.html to server
- If the url contained a different path, it will appear in the request line

GET / HTTP/1.1

Host: cse312.com

GET /lecture HTTP/1.1

Host: cse312.com

Headers

- Following the request line are any number of headers
- HTTP Headers
 - Key-Value pairs
 - Key and value separated by a colon ":"
- Each header will be a new line
- To parse, look for the colon ":" and read the key and value

GET / HTTP/1.1

Host: cse312.com

Response

- Your web server will listen for requests over the TCP sockets and respond with HTTP responses
- Send this response back to the client to serve them the requested content

HTTP/1.1 200 OK

Content-Type: text/plain

Content-Length: 5

hello

Status Line

- The first line of the response must be the status line
- Status line contains 3 values separated by spaces
 - The HTTP version
 - The status code
 - The status message (Reason phrase in docs)

HTTP/1.1 200 OK

Content-Type: text/plain

Content-Length: 5

hello

Response Headers

- The headers in the response follow the same format as request headers
- Should have at least two headers
 - Content-Type - Tells the browser how to parse this content
 - Content-Length - How much information should be read from the body

HTTP/1.1 200 OK

Content-Type: text/plain
Content-Length: 5

hello

Body

- The headers are followed by 2 new lines to indicate the beginning of the body
- The body contains the content that is being served
- The body ends with 2 new line characters
- [Requests can have bodies too]

HTTP/1.1 200 OK

Content-Type: text/plain

Content-Length: 5

hello

Response Codes

- Tells the browser the nature of the response
 - 200 means everything went well
 - 404 means the content was not found
- Include a human readable message
 - OK
 - Each code has a standard message, but different messages are accepted

HTTP/1.1 200 OK

Content-Type: text/plain

Content-Length: 5

hello

Response Codes

- Some responses don't include a body
- To redirect to another page use 301 with no body
 - Must contain a Location header for the location of the redirect
 - Location can be a relative path (Same server) or an absolute path

HTTP/1.1 301 OK
Location: /