

Databases

Databases

- Software that stores data on disk
- Runs as a server and is communicated with via TCP sockets
- Provides an API to store/retrieve data
 - The software handles the low-level file IO
 - Allows us to think about our data, not how to store it
- Provides many optimizations

Databases

- We'll look at 2 different database
- Both are pieces of software that must be downloaded, installed, ran, then connected to via TCP
- MySQL
 - A server implementing SQL (Structured Query Language)
- MongoDB
 - A server based on document stores

MySQL

- One you download, install, and run the server
 - It will listen for TCP connections on port 3306 (By default)
- Install a library for your language that will connect to the MySQL server
 - You will not have to connect to your database at the TCP level in this course (True for MongoDB as well)
 - The library will provide a convenient API
 - Send queries using the query language

MySQL

- After MySQL is running and you install a library to connect to it
- Connect to MySQL Server by providing
 - The url of the database
 - username/password for the database
 - Whatever you chose when setting up the database

```
val url = "jdbc:mysql://localhost/mysql"  
val username = "root"  
val password = "12345678"
```

```
var connection: Connection = DriverManager.getConnection(url, username, password)
```

MySQL - Security

- For real apps that you deploy
 - **Do not check your password into version control!**
 - A plain text password in public GitHub repo is bad
 - Attacker can replace localhost with the IP for your app and can access all your data
 - Common to save the password in a environment variable to prevent accidentally pushing it to git
 - **Do not use the default password for any servers you're running**
 - This is what caused the Equifax leak (Not with MySQL)
- Attacker have bots that scan random IPs for such vulnerabilities

```
val url = "jdbc:mysql://localhost/mysql?serverTimezone=UTC"
val username = "root"
val password = "12345678"
```

```
var connection: Connection = DriverManager.getConnection(url, username, password)
```

MySQL - Security

- Can use Docker to set an environment variable containing your DB password
 - Do not add the password when checking it into the repo
- When you're ready to deploy the app
 - Clone the repo, choose a password, and edit the files on the production server only
 - Access to this password should be on a need-to-know basis
- Alternatively/Additionally: Change your DB settings to only allow connections from localhost (Unless your app is distributed)

```
val url = "jdbc:mysql://localhost/mysql?serverTimezone=UTC"
val username = "root"
val password = "12345678"
```

```
var connection: Connection = DriverManager.getConnection(url, username, password)
```

MySQL

- Once connected, we can send SQL statements to the server

```
val statement = connection.createStatement()  
statement.execute("CREATE TABLE IF NOT EXISTS players (username TEXT, points INT)")
```

- If using inputs from the user always use prepared statements
- Indices start at 1 in this example 🥲

```
val statement = connection.prepareStatement("INSERT INTO players VALUE (?, ?)")  
  
statement.setString(1, "mario")  
statement.setInt(2, 10)  
  
statement.execute()
```


MySQL - Security

- Not using prepared statements?
 - **Vulnerable to SQL injection attacks**
- If you concatenate user inputs directly into your SQL statements
 - Attacker chooses a username of `"";DROP TABLE players;`
 - You lose all your data
 - Even worse, they find a way to access the entire database and steal other users' data
 - SQL Injection is the most common successful attack on servers

MySQL

- Send queries to pull data from the database
- Returns a ResultSet in this example
 - The next() methods queue the next result of the query
 - next returns false if there are no more results to read
- Can read values by index or by column name

```
val statement = connection.createStatement()
val result: ResultSet = statement.executeQuery("SELECT * FROM players")

var allScores: Map[String, Int] = Map()

while (result.next()) {
    val username = result.getString("username")
    val score = result.getInt("points")
    allScores = allScores + (username -> score)
}
```

SQL

- SQL is based on tables with rows and column
 - Similar in structure to CSV except the values have types other than string
- How do we store an array or key-value store?
 - With CSV our answer was to move on to JSON
 - SQL answer is to create a separate table and use JOINS
 - Or, try MongoDB

MongoDB

- Runs on port 27017 (By default)
- A document-based database
- Instead of using tables, stores data in a structure very similar to JSON
- In python/JS
 - Insert dictionaries/objects directly
- Each object is stored in a collection

```
chat_collection.insert_one({'username': 'hartloff', 'message': 'hello'})
```

MongoDB

- Retrieve documents using find
- Find takes a key-value store and returns all documents with those values stored at the given keys
 - Ex. {'username': 'hartloff'} returns all documents with a username of hartloff
- To retrieve all documents, use an empty key-value store {}

```
collection.find({'username': 'hartloff'})  
collection.find({})
```

MongoDB vs. SQL

- MongoDB is unstructured
 - Can add objects in any format to a collection
 - Can mix formats in a single collection
 - I.e. In a single collection the documents can have different attributes
- SQL is structured (That's what the S stands for)
 - Table columns must be pre-defined
 - All rows have the same attributes
 - Adding a column can be difficult
 - Fast!

MongoDB vs. SQL

- Hot Take
 - MongoDB is best for prototyping when the structure of your data is constantly changing
 - Take advantage of the flexibility
 - SQL is best once your data has a defined structure
 - Take advantage of the efficiency